

The SOLID Principles

The SOLID Principles are five principles of Object-Oriented class design. They are a set of rules and best practices to follow while designing a class structure.

SOLID stands for:

S - Single-responsibility Principle

O - Open-closed Principle

L - Liskov Substitution Principle

I - Interface Segregation Principle

D - Dependency Inversion Principle

Single-Responsibility Principle (SRP)

The **single-responsibility principle (SRP)** comes from [Robert C. Martin](#), more commonly known by his nickname Uncle Bob, who's a well-respected figure in the software engineering world and one of the original signatories of the [Agile Manifesto](#). In fact, he coined the term SOLID.

The single-responsibility principle states that:

A class should have only one reason to change.

This means that a class should have only one **responsibility**, as expressed through its methods. If a class takes care of more than one task, then you should separate those tasks into separate classes.

This principle is closely related to the concept of [separation of concerns](#), which suggests that you should split your programs into different sections. Each section must address a separate concern.

To illustrate the single-responsibility principle and how it can help you improve your object-oriented design, say that you have the following FileManager class:

```
# file_manager_srp.py

from pathlib import Path
from zipfile import ZipFile

class FileManager:
    def __init__(self, filename):
        self.path = Path(filename)

    def read(self, encoding="utf-8"):
        return self.path.read_text(encoding)

    def write(self, data, encoding="utf-8"):
        self.path.write_text(data, encoding)

    def compress(self):
        with ZipFile(self.path.with_suffix(".zip"), mode="w") as archive:
            archive.write(self.path)

    def decompress(self):
        with ZipFile(self.path.with_suffix(".zip"), mode="r") as archive:
            archive.extractall()
```

In this example, your `FileManager` class has two different responsibilities. It uses the `.read()` and `.write()` methods to manage the file. It also deals with [ZIP archives](#) by providing the `.compress()` and `.decompress()` methods.

This class violates the single-responsibility principle because it has two reasons for changing its internal implementation. To fix this issue and make your design more robust, you can split the class into two smaller, more focused classes, each with its own specific concern:

```
# file_manager_srp.py

from pathlib import Path
from zipfile import ZipFile

class FileManager:
    def __init__(self, filename):
```

```

        self.path = Path(filename)

    def read(self, encoding="utf-8"):
        return self.path.read_text(encoding)

    def write(self, data, encoding="utf-8"):
        self.path.write_text(data, encoding)

class ZipFileManager:
    def __init__(self, filename):
        self.path = Path(filename)

    def compress(self):
        with ZipFile(self.path.with_suffix(".zip"), mode="w") as archive:
            archive.write(self.path)

    def decompress(self):
        with ZipFile(self.path.with_suffix(".zip"), mode="r") as archive:
            archive.extractall()

```

Now you have two smaller classes, each having only a single responsibility. `FileManager` takes care of managing a file, while `ZipFileManager` handles the [compression](#) and [decompression](#) of a file using the ZIP format. These two classes are smaller, so they're more manageable. They're also easier to reason about, test, and debug.

The concept of **responsibility** in this context may be pretty subjective. Having a single responsibility doesn't necessarily mean having a single method. Responsibility isn't directly tied to the number of methods but to the core task that your class is responsible for, depending on your idea of what the class represents in your code. However, that subjectivity shouldn't stop you from striving to use the SRP.

Open-Closed Principle (OCP)

The **open-closed principle (OCP)** for object-oriented design was originally introduced by [Bertrand Meyer](#) in 1988 and means that:

Software entities (classes, modules, functions, etc.) should be open for extension, but closed for modification.

To understand what the open-closed principle is all about, consider the following `Shape` class:

```
# shapes_ocp.py

from math import pi

class Shape:
    def __init__(self, shape_type, **kwargs):
        self.shape_type = shape_type
        if self.shape_type == "rectangle":
            self.width = kwargs["width"]
            self.height = kwargs["height"]
        elif self.shape_type == "circle":
            self.radius = kwargs["radius"]

    def calculate_area(self):
        if self.shape_type == "rectangle":
            return self.width * self.height
        elif self.shape_type == "circle":
            return pi * self.radius**2
```

The initializer of Shape takes a `shape_type` argument that can be either "rectangle" or "circle". It also takes a specific set of keyword arguments using the `**kwargs` syntax. If you set the shape type to "rectangle", then you should also pass the width and height keyword arguments so that you can construct a proper rectangle.

In contrast, if you set the shape type to "circle", then you must also pass a radius argument to construct a circle.

Shape also has a `.calculate_area()` method that computes the area of the current shape according to its `.shape_type`:

```
>>>
>>> from shapes_ocp import Shape

>>> rectangle = Shape("rectangle", width=10, height=5)
>>> rectangle.calculate_area()
50
>>> circle = Shape("circle", radius=5)
>>> circle.calculate_area()
78.53981633974483
```

The class works. You can create circles and rectangles, compute their area, and so on. However, the class looks pretty bad. Something seems wrong with it at first sight.

Imagine that you need to add a new shape, maybe a square. How would you do that? Well, the option here is to add another `elif` clause to `.__init__()` and to `.calculate_area()` so that you can address the requirements of a square shape.

Having to make these changes to create new shapes means that your class is open to modification. That violates the open-closed principle. How can you fix your class to make it open to extension but closed to modification? Here's a possible solution:

```
# shapes_ocp.py

from abc import ABC, abstractmethod
from math import pi

class Shape(ABC):
    def __init__(self, shape_type):
        self.shape_type = shape_type

    @abstractmethod
    def calculate_area(self):
        pass

class Circle(Shape):
    def __init__(self, radius):
        super().__init__("circle")
        self.radius = radius

    def calculate_area(self):
        return pi * self.radius**2

class Rectangle(Shape):
    def __init__(self, width, height):
        super().__init__("rectangle")
        self.width = width
        self.height = height
```

```

    def calculate_area(self):
        return self.width * self.height

class Square(Shape):
    def __init__(self, side):
        super().__init__("square")
        self.side = side

    def calculate_area(self):
        return self.side**2

```

In this code, you completely [refactored](#) the Shape class, turning it into an [abstract base class \(ABC\)](#). This class provides the required [interface \(API\)](#) for any shape that you'd like to define. That [interface](#) consists of a `.shape_type` attribute and a `.calculate_area()` method that you must override in all the subclasses.

This update closes the class to modifications. Now you can add new shapes to your class design without the need to modify Shape. In every case, you'll have to implement the required interface, which also makes your classes [polymorphic](#).

Liskov Substitution Principle (LSP)

The **Liskov substitution principle (LSP)** was introduced by [Barbara Liskov](#) at an [OOPSLA conference](#) in 1987. Since then, this principle has been a fundamental part of object-oriented programming. The principle states that:

Subtypes must be substitutable for their base types.

For example, if you have a piece of code that works with a Shape class, then you should be able to substitute that class with any of its subclasses, such as Circle or Rectangle, without breaking the code.

In practice, this principle is about making your subclasses behave like their base classes without breaking anyone's expectations when they call the same methods. To continue with shape-related examples, say you have a Rectangle class like the following:

```

# shapes_lsp.py

class Rectangle:
    def __init__(self, width, height):
        self.width = width
        self.height = height

```

```
def calculate_area(self):  
    return self.width * self.height
```

In Rectangle, you've provided the `.calculate_area()` method, which operates with the `.width` and `.height` [instance attributes](#).

Because a square is a special case of a rectangle with equal sides, you think of deriving a Square class from Rectangle in order to reuse the code. Then, you override the [setter](#) method for the `.width` and `.height` attributes so that when one side changes, the other side also changes:

```
# shapes_lsp.py  
  
# ...  
  
class Square(Rectangle):  
    def __init__(self, side):  
        super().__init__(side, side)  
  
    def __setattr__(self, key, value):  
        super().__setattr__(key, value)  
        if key in ("width", "height"):  
            self.__dict__["width"] = value  
            self.__dict__["height"] = value
```

In this snippet of code, you've defined Square as a subclass of Rectangle. As a user might expect, the class [constructor](#) takes only the side of the square as an argument. Internally, the `__init__()` method initializes the parent's attributes, `.width` and `.height`, with the side argument.

You've also defined a [special method](#), `__setattr__()`, to hook into Python's attribute-setting mechanism and intercept the [assignment](#) of a new value to either the `.width` or `.height` attribute. Specifically, when you set one of those attributes, the other attribute is also set to the same value:

```
>>>  
>>> from shapes_lsp import Square  
  
>>> square = Square(5)  
>>> vars(square)  
{'width': 5, 'height': 5}
```

```
>>> square.width = 7
>>> vars(square)
{'width': 7, 'height': 7}
```

```
>>> square.height = 9
>>> vars(square)
{'width': 9, 'height': 9}
```

Now you've ensured that the Square object always remains a valid square, making your life easier for the small price of a bit of wasted memory. Unfortunately, this violates the Liskov substitution principle because you can't replace instances of Rectangle with their Square counterparts.

When someone expects a rectangle object in their code, they might assume that it'll behave like one by exposing two independent `.width` and `.height` attributes. Meanwhile, your Square class breaks that assumption by changing the behavior promised by the object's interface. That could have surprising and unwanted consequences, which would likely be hard to [debug](#).

While a square is a specific type of rectangle in mathematics, the classes that represent those shapes shouldn't be in a parent-child relationship if you want them to comply with the Liskov substitution principle. One way to solve this problem is to create a base class for both Rectangle and Square to extend:

```
# shapes_lsp.py

from abc import ABC, abstractmethod

class Shape(ABC):
    @abstractmethod
    def calculate_area(self):
        pass

class Rectangle(Shape):
    def __init__(self, width, height):
        self.width = width
        self.height = height

    def calculate_area(self):
        return self.width * self.height
```



```
class Square(Shape):
    def __init__(self, side):
        self.side = side

    def calculate_area(self):
        return self.side ** 2
```

Shape becomes the type that you can substitute through polymorphism with either Rectangle or Square, which are now siblings rather than a parent and a child. Notice that both concrete shape types have distinct sets of attributes, different initializer methods, and could potentially implement even more separate behaviors. The only thing that they have in common is the ability to calculate their area.

With this implementation in place, you can use the Shape type interchangeably with its Square and Rectangle subtypes when you only care about their common behavior:

```
>>>
>>> from shapes_lsp import Rectangle, Square

>>> def get_total_area(shapes):
...     return sum(shape.calculate_area() for shape in shapes)

>>> get_total_area([Rectangle(10, 5), Square(5)])
75
```

Here, you pass a pair consisting of a rectangle and a square into a function that calculates their total area. Because the function only cares about the `.calculate_area()` method, it doesn't matter that the shapes are different. This is the essence of the Liskov substitution principle.

Interface Segregation Principle (ISP)

The **interface segregation principle (ISP)** comes from the same mind as the single-responsibility principle. Yes, it's another feather in [Uncle Bob's](#) cap. The principle's main idea is that:

Clients should not be forced to depend upon methods that they do not use. Interfaces belong to clients, not to hierarchies.

In this case, *clients* are classes and subclasses, and *interfaces* consist of methods and attributes. In other words, if a class doesn't use particular methods or attributes, then those methods and attributes should be segregated into more specific classes.

Consider the following example of class [hierarchy](#) to model printing machines:

```
# printers_isp.py

from abc import ABC, abstractmethod

class Printer(ABC):
    @abstractmethod
    def print(self, document):
        pass

    @abstractmethod
    def fax(self, document):
        pass

    @abstractmethod
    def scan(self, document):
        pass

class OldPrinter(Printer):
    def print(self, document):
        print(f"Printing {document} in black and white...")

    def fax(self, document):
        raise NotImplementedError("Fax functionality not supported")

    def scan(self, document):
        raise NotImplementedError("Scan functionality not supported")

class ModernPrinter(Printer):
    def print(self, document):
        print(f"Printing {document} in color...")

    def fax(self, document):
        print(f"Faxing {document}...")
```

```
def scan(self, document):  
    print(f"Scanning {document}...")
```

In this example, the base class, `Printer`, provides the interface that its subclasses must implement. `OldPrinter` inherits from `Printer` and must implement the same interface. However, `OldPrinter` doesn't use the `.fax()` and `.scan()` methods because this type of printer doesn't support these functionalities.

This implementation violates the ISP because it forces `OldPrinter` to expose an interface that the class doesn't implement or need. To fix this issue, you should separate the interfaces into smaller and more specific classes. Then you can create concrete classes by inheriting from multiple interface classes as needed:

```
# printers_isp.py  
  
from abc import ABC, abstractmethod  
  
class Printer(ABC):  
    @abstractmethod  
    def print(self, document):  
        pass  
  
class Fax(ABC):  
    @abstractmethod  
    def fax(self, document):  
        pass  
  
class Scanner(ABC):  
    @abstractmethod  
    def scan(self, document):  
        pass  
  
class OldPrinter(Printer):  
    def print(self, document):  
        print(f"Printing {document} in black and white...")  
  
class NewPrinter(Printer, Fax, Scanner):  
    def print(self, document):  
        print(f"Printing {document} in color...")
```

```
def fax(self, document):
    print(f"Faxing {document}...")

def scan(self, document):
    print(f"Scanning {document}...")
```

Now Printer, Fax, and Scanner are base classes that provide specific interfaces with a single responsibility each. To create OldPrinter, you only inherit the Printer interface. This way, the class won't have unused methods. To create the ModernPrinter class, you need to inherit from all the interfaces. In short, you've segregated the Printer interface.

This class design allows you to create different machines with different sets of functionalities, making your design more flexible and extensible.

Dependency Inversion Principle (DIP)

The **dependency inversion principle (DIP)** is the last principle in the SOLID set. This principle states that:

Abstractions should not depend upon details. Details should depend upon abstractions.

That sounds pretty complex. Here's an example that will help to clarify it. Say you're building an application and have a FrontEnd class to display data to the users in a friendly way. The app currently gets its data from a database, so you end up with the following code:

```
# app_dip.py

class FrontEnd:
    def __init__(self, back_end):
        self.back_end = back_end

    def display_data(self):
        data = self.back_end.get_data_from_database()
        print("Display data:", data)

class BackEnd:
```

```
def get_data_from_database(self):  
    return "Data from the database"
```

In this example, the FrontEnd class depends on the BackEnd class and its concrete implementation. You can say that both classes are tightly coupled. This coupling can lead to scalability issues. For example, say that your app is growing fast, and you want the app to be able to read data from a [REST API](#). How would you do that?

You may think of adding a new method to BackEnd to retrieve the data from the REST API. However, that will also require you to modify FrontEnd, which should be closed to modification, according to the [open-closed principle](#).

To fix the issue, you can apply the dependency inversion principle and make your classes depend on abstractions rather than on concrete implementations like BackEnd. In this specific example, you can introduce a DataSource class that provides the interface to use in your concrete classes:

```
# app_dip.py  
  
from abc import ABC, abstractmethod  
  
class FrontEnd:  
    def __init__(self, data_source):  
        self.data_source = data_source  
  
    def display_data(self):  
        data = self.data_source.get_data()  
        print("Display data:", data)  
  
class DataSource(ABC):  
    @abstractmethod  
    def get_data(self):  
        pass  
  
class Database(DataSource):  
    def get_data(self):
```

```
        return "Data from the database"

class API(DataSource):
    def get_data(self):
        return "Data from the API"
```

In this redesign of your classes, you've added a `DataSource` class as an abstraction that provides the required interface, or the `.get_data()` method. Note how `FrontEnd` now depends on the interface provided by `DataSource`, which is an abstraction.

Then you define the `Database` class, which is a concrete implementation for those cases where you want to retrieve the data from your database. This class depends on the `DataSource` abstraction through inheritance. Finally, you define the `API` class to support retrieving the data from the REST API. This class also depends on the `DataSource` abstraction.

Here's how you can use the `FrontEnd` class in your code:

```
>>>
>>> from app_dip import API, Database, FrontEnd

>>> db_front_end = FrontEnd(Database())
>>> db_front_end.display_data()
Display data: Data from the database

>>> api_front_end = FrontEnd(API())
>>> api_front_end.display_data()
Display data: Data from the API
```

Here, you first initialize `FrontEnd` using a `Database` object and then again using an `API` object. Every time you call `.display_data()`, the result will depend on the concrete data source that you use. Note that you can also change the data source dynamically by reassigning the `.data_source` attribute in your `FrontEnd` instance.